

Baixiang Huang

baixianghuang.github.io [Google Scholar](#) [LinkedIn](#) baixiang.huang@emory.edu

EDUCATION

Emory University	Jan 2025 – Present
Ph.D. in Computer Science	Advisor: Kai Shu
Illinois Institute of Technology	Jan 2023 – Dec 2024
Ph.D. in Computer Science (Transfer Out)	Advisor: Kai Shu
National University of Singapore	Aug 2020 – Feb 2022
Master’s Degree in Computer Science	Advisor: Bryan Hooi
Central South University	Sep 2015 – Jun 2019
Bachelor’s Degree in Computer Science	

PROFESSIONAL EXPERIENCE

Amazon Search	May 2025 – Dec 2025
Applied Scientist Intern	Mentor: Limeng Cui
Proposed a model editing framework for LLM personalization; paper currently under review.	

PUBLICATIONS AND PREPRINTS

Peer-Reviewed Conference and Journal Publications

- [1] Can Knowledge Editing Really Correct Hallucinations?
Baixiang Huang, Canyu Chen, Xiong Xiao Xu, Ali Payani, Kai Shu.
The Thirteenth International Conference on Learning Representations (ICLR). 2025.
[PDF] [GitHub] [Project Website]
- [2] SST: Multi-Scale Hybrid Mamba-Transformer Experts for Long-Short Range Time Series Forecasting
Xiong Xiao Xu, Canyu Chen, Yueqing Liang, **Baixiang Huang**, Guangji Bai, Liang Zhao, Kai Shu.
Proceedings of the 34th ACM International Conference on Information and Knowledge Management (CIKM), 2025. [PDF] [GitHub]
- [3] Can Large Language Models Identify Authorship?
Baixiang Huang, Canyu Chen, Kai Shu.
Empirical Methods in Natural Language Processing (EMNLP Findings). 2024.
[PDF] [GitHub] [Project Website]
- [4] Authorship Attribution in the Era of LLMs: Problems, Methodologies, and Challenges
Baixiang Huang, Canyu Chen, Kai Shu.
ACM SIGKDD Explorations. 2024. [PDF] [GitHub] [Project Website]
- [5] TAP: A Comprehensive Data Repository for Traffic Accident Prediction in Road Networks
Baixiang Huang, Bryan Hooi, Kai Shu.
Proceedings of the 31st ACM International Conference on Advances in Geographic Information Systems (ACM SIGSPATIAL). 2023. [PDF] [GitHub]

Preprints

- [1] Model Editing as a Double-Edged Sword: Steering Agent Ethical Behavior Toward Beneficence or Harm
Baixiang Huang, Zhen Tan, Haoran Wang, Zijie Liu, Dawei Li, Ali Payani, Huan Liu, Tianlong Chen, and Kai Shu.
arXiv preprint 2025. [PDF] [GitHub] [Project Website]
- [2] Privacy-Aware Decoding: Mitigating Privacy Leakage of Large Language Models in Retrieval-Augmented Generation
Haoran Wang, Xiong Xiao Xu, **Baixiang Huang**, and Kai Shu.
arXiv preprint 2025. [PDF] [GitHub]
- [3] Can Editing LLMs Inject Harm?
Canyu Chen*, **Baixiang Huang***, Zekun Li, Zhaorun Chen, Shiyang Lai, Xiong Xiao Xu, Jia-Chen Gu, Jindong Gu, Huaxiu Yao, Chaowei Xiao, Xifeng Yan, William Wang, Philip Torr, Dawn Song, Kai Shu.
arXiv preprint 2024. [PDF] [GitHub] [Project Website]

* denotes equal contribution

PROJECT EXPERIENCE

The Human Interpretable Attribution of Text Using Underlying Structure Jan 2023 – May 2024
Sponsor: The Intelligence Advanced Research Projects Activity (IARPA) Advisor: Kai Shu

- Developed a contrastive learning framework based on DeBERTa to model authorship representations, enabling differentiation of texts by distinct authors.
- Built and evaluated a large-scale cross-genre Reddit dataset with hard negative samples.
- Improved the cross-domain performance of the authorship representation learning framework using new data.
- Proposed a linguistically guided technique leveraging Large Language Models for zero-shot authorship verification and attribution.

TEACHING

- Undergraduate TA for CS253: Data Structures and Algorithms, Fall 2025
- Undergraduate TA for CS470: Data Mining, Spring 2025
- Graduate TA for CS536: Science of Programming, Fall 2024

ACADEMIC SERVICE

- Reviewer: ICLR {2024, 2025}, KDD {2024}, SIGIR {2023, 2024}, CIKM {2023}, ICDM {2023}, TheWebConf {2023}
- Program Committee: AAI {2023, 2024, 2025}

SKILLS

- Frameworks: PyTorch, TensorFlow, MyBatis, Spring Boot, Vue
- Miscellaneous: Server Hardware Maintenance, L^AT_EX, SQL, Git